

Exercises - Yang

These exercises will make you feel comfortable with YANG. Parts 1, 2 and 4 are based on the ONF Connect 2019 Next-Gen SDN Tutorial. Source material can be obtained: <https://github.com/opennetworkinglab/ngsdn-tutorial>.

This exercise is divided in three parts (and an extra 4th):

1. Understanding the YANG language
2. Understand YANG encoding
3. Understanding YANG-enabled API
4. Understanding YANG-enabled transport protocols

Part 1: Understanding the YANG language

We will start with a simple YANG module called `demo-port` in `yang/demo-port.yang`

Take a look at the model and try to derive the structure. What are valid values for each of the leaf nodes?

This model is self contained, so it isn't too difficult to work it out. However, most YANG models are defined over many files that makes it very complicated to work out the overall structure.

To make this easier, we can use a tool called `pyang` to try to visualize the structure of the model.

Open a terminal, enter the `yang-tutorial` folder, and run `pyang` on the `demo-port.yang` model:

```
$ cd yang-tutorial
$ pyang -f tree demo-port.yang
```

You should see a tree representation of the `demo-port` module. Does this match your expectations?

Extra Credit: If you finish this quickly, you can try to add a new leaf node to `port-config` or `port-state` grouping, then rerun `pyang` and see where your new leaf was added.

We can also use `pyang` to visualize a more complicated set of models, like the set of OpenConfig models that Stratum uses.

These models have already been loaded into the `yang-tools` container in the `models` directory.

```
$ cd models
$ pyang -f tree \
  -p ietf \
  -p openconfig \
```

```

    -p hercules \
    openconfig/interfaces/openconfig-interfaces.yang \
    openconfig/interfaces/openconfig-if-ethernet.yang \
    openconfig/platform/* \
    openconfig/qos/* \
    openconfig/system/openconfig-system.yang \
    hercules/openconfig-hercules-*.yang | less
$ cd -

```

You should see a tree structure of the models displayed in `less`. You can use the Arrow keys or j/k to scroll up and down. Type q to quit.

In the interface model, we can see the path to enable or disable an interface: `interfaces/interface[name]/config/enabled`

What is the path to read the number of incoming packets (`in-pkts`) on an interface?

Extra Credit: If you have some time, take a look at the models in the `/models` directory or browse them on Github: <https://github.com/openconfig/public/tree/master/release/models>

Try to find the description of the `enabled` or `in-pkts` leaf nodes.

Hint: Take a look at the `openconfig-interfaces.yang` file.

Part 2: Understand YANG encoding

There is no specific YANG data encoding, but data adhering to YANG models can be encoded into XML, JSON, or Protobuf (among other formats). Each of these formats has it's own schema format.

First, we can look at YANG's first and canonical representation format XML. To see a empty skeleton of data encoded in XML, run:

```
$ pyang -f sample-xml-skeleton demo-port.yang
```

This skeleton should match the tree representation we saw in part 1.

We can also use `pyang` to generate a DSDL schema based on the YANG model:

```
$ pyang -f dsdl demo-port.yang | xmllint --format -
```

The first part of the schema describes the tree structure, and the second part describes the value constraints for the leaf nodes.

Extra credit: Try adding new speed identity (e.g. `SPEED_100G`) or changing the range for `port-number` values in `demo-port.yang`, then rerun `pyang -f dsdl`. Do you see your changes reflected in the DSDL schema?

Next, we will look at encoding data using Protocol Buffers (protobuf). The protobuf encoding is a more compact binary encoding than XML, and libraries can be automatically generated for dozens of languages. We can use `ygot`'s `proto_generator` to generate protobuf messages from our YANG model.

More info about `ygot` see `Getting Started`.

```
$ protogenerator -output_dir=./proto -package_name=tutorial demo-port.yang
protogenerator will generate two files: * /proto/tutorial/demo_port/demo_port.proto
* /proto/tutorial/enums/enums.proto
```

Open `demo_port.proto` using `less`:

```
$ less proto/tutorial/demo_port/demo_port.proto
```

This file contains a top-level `Ports` message that matches the structure defined in the YANG model. You can see that `protogenerator` also adds a `yext.schemapath` custom option to each protobuf message field that explicitly maps to the YANG leaf path. Enums (like `tutorial.enums.DemoPortSPEED`) aren't included in this file, but `protogenerator` puts them in a separate file: `enums.proto`

Open `enums.proto` using `less`:

```
$ less proto/tutorial/enums/enums.proto
```

You should see an enum for the 10GB speed, along with any other speeds that you added if you completed the extra credit above.

We can also use `protogenerator` to build the protobuf messages for the Open-Config models that Stratum uses:

```
$ cd models
```

```
$ protogenerator \
  -generate_fakeroot \
  -output_dir=./proto \
  -package_name=openconfig \
  -exclude_modules=ietf-interfaces \
  -compress_paths \
  -base_import_path= \
  -path=ietf,openconfig,hercules \
  openconfig/interfaces/openconfig-interfaces.yang \
  openconfig/interfaces/openconfig-if-ip.yang \
  openconfig/lacp/openconfig-lacp.yang \
  openconfig/platform/openconfig-platform-linecard.yang \
  openconfig/platform/openconfig-platform-port.yang \
  openconfig/platform/openconfig-platform-transceiver.yang \
  openconfig/platform/openconfig-platform.yang \
```

```
openconfig/system/openconfig-system.yang \  
openconfig/vlan/openconfig-vlan.yang \  
hercules/openconfig-hercules-interfaces.yang \  
hercules/openconfig-hercules-platform-chassis.yang \  
hercules/openconfig-hercules-platform-linecard.yang \  
hercules/openconfig-hercules-platform-node.yang \  
hercules/openconfig-hercules-platform-port.yang \  
hercules/openconfig-hercules-platform.yang \  
hercules/openconfig-hercules-qos.yang \  
hercules/openconfig-hercules.yang
```

```
$ cd -
```

You will find `openconfig.proto` and `enums.proto` in the `proto/openconfig` directory.

Extra Credit: Try to find the Protobuf message fields used to enable a port or get the ingress packets counter in the protobuf messages.

Hint: Searching by `schemapath` might help.

`ygot` can also be used to generate Go structs that adhere to the YANG model and that are capable of validating the structure, type, and values of data.

Extra Credit: If you have extra time or are interested in using YANG and Go together, try generating Go code for the `demo-port` module.

```
$ mkdir -p ./goSrc  
$ generator -output_dir=./goSrc -package_name=tutorial demo-port.yang
```

Take a look at the Go files in `/goSrc`.

You can now quit out of the container (using `Ctrl-D` or `exit`).

Part 3: Understanding YANG-enabled API

This part used `pyang` to generate a python API that is used to fill the YANG model and explore parsing and serialization in different formats.

Use `pyang` and `pyandbind` to generate all the doc files for `demo-port.yang` and validate the file in `examples/` folder against the python generated source code `demo_port.py`.

```
$ make all
```

Check in the `./doc` folder the created `.uml` `.tree` and `.html` files.

Open the source code file named `demo_port.py` and see the syntax generated by `pyangbind`.

And finally take a look at the file located in the examples folder named `demo_ports_01.yaml`. See how this structure matches the tree elaborated by the doc file.

Now, check the construction of a python library from scratch, and how it can be serialized into JSON.

```
$ /usr/bin/python3 serialize_demo_port.py
```

Open the file `serialize_demo_port.py` and check how to load the `demo_port.py` source file and fill the data structure offered by the YANG module `demo-port`.

Part 4 (EXTRA): Understanding YANG-enabled transport protocols

This is the last part of the: ONF Connect 2019 Next-Gen SDN Tutorial. Source material can be obtained: <https://github.com/opennetworkinglab/ngsdn-tutorial>.

To build the dependencies for this part, follow the steps below.

```
$ git clone https://github.com/opennetworkinglab/ngsdn-tutorial
$ cd ~/ngsdn-tutorial
$ git pull origin master
$ make pull-deps
```

There are several YANG-model agnostic protocols that can be used to get or set data that adheres to a model, like NETCONF, RESTCONF, and gNMI.

This part focuses on using the protobuf encoding over gNMI.

First, make sure that your Mininet container is still running.

```
$ make start
docker-compose up -d
mininet is up-to-date
onos is up-to-date
```

If you see the following output, then Mininet was not running:

```
Starting mininet ... done
Starting onos    ... done
```

You will need to go back to Exercise 1 and install forwarding rules to re-establish pings between `h1a` and `h1b` for later parts of this exercise.

Next, we will use a gNMI client CLI to read the all of the configuration from the Stratum switch `leaf1` in our Mininet network:

```
$ util/gnmi-cli --grpc-addr localhost:50001 get /
```

The first part of the output shows the request that was made by the CLI:

```
REQUEST
path {
}
type: CONFIG
encoding: PROTO
```

The path being requested is the empty path (which means the root of the config tree), the type of data is just the config tree, and the requested encoding for the response is protobuf.

The second part of the output shows the response from Stratum:

```
RESPONSE
notification {
  update {
    path {
    }
    val {
      any_val {
        type_url: "type.googleapis.com/openconfig.Device"
        value: \252\221\231\304\001\... TRUNCATED
      }
    }
  }
}
```

You can see that Stratum provides a response of type `openconfig.Device`, which is the top-level message defined in `openconfig.proto`. The response is the binary encoding of the data based on the protobuf message.

The value is not very human readable, but we can translate the reply using a utility that converts between the binary and textual representations of the protobuf message.

We can rerun the command, but this time pipe the output through the converter utility (then pipe that output to `less` to make scrolling easier):

```
$ util/gnmi-cli --grpc-addr localhost:50001 get / | util/oc-pb-decoder | less
```

The contents of the response should be easier to read. Scroll down to the first **interface**. Is the interface enabled? What is the speed of the port?

Extra credit: Can you find **in-pkts**? If not, why do you think they are missing?

One of the benefits to gNMI is it's "schema-less" encoding that allows clients or devices to update only the paths that need to be updated. This is particularly useful for subscriptions.

First, let's try out the schema-less representation by requesting the configuration port between `leaf1` and `h1a`:

```
$ util/gnmi-cli --grpc-addr localhost:50001 get \  
  /interfaces/interface[name=leaf1-eth3]/config
```

You should see this response containing 2 leafs under config - **enabled** and **health-indicator**:

```
RESPONSE  
notification {  
  update {  
    path {  
      elem {  
        name: "interfaces"  
      }  
      elem {  
        name: "interface"  
        key {  
          key: "name"  
          value: "leaf1-eth3"  
        }  
      }  
    }  
    elem {  
      name: "config"  
    }  
    elem {  
      name: "enabled"  
    }  
  }  
  val {  
    bool_val: true  
  }  
}  
notification {  
  update {  
    path {  
      elem {  
        name: "interfaces"  
      }  
      elem {  
        name: "interface"  
        key {  
          key: "name"  
          value: "leaf1-eth3"  
        }  
      }  
    }  
    elem {  
      name: "config"
```

```

    }
    elem {
      name: "health-indicator"
    }
  }
  val {
    string_val: "GOOD"
  }
}
}

```

The schema-less representation provides an **update** for each leaf containing both the path and the value of the leaf. You can confirm that the interface is enabled (set to **true**).

Next, we will subscribe to the ingress unicast packet counters for the interface on leaf1 attached to h1a (port 3):

```

$ util/gnmi-cli --grpc-addr localhost:50001 \
  --interval 1000 sub-sample \
  /interfaces/interface[name=leaf1-eth3]/state/counters/in-unicast-pkts

```

The first part of the output shows the request being made by the CLI:

```

REQUEST
subscribe {
  subscription {
    path {
      elem {
        name: "interfaces"
      }
      elem {
        name: "interface"
        key {
          key: "name"
          value: "leaf1-eth3"
        }
      }
    }
    elem {
      name: "state"
    }
    elem {
      name: "counters"
    }
    elem {
      name: "in-unicast-pkts"
    }
  }
}

```



```

    mode: SAMPLE
    sample_interval: 1000
  }
  updates_only: true
}

```

We have the subscription path, the type of subscription (sampling) and the sampling rate (every 1000ms, or 1s).

The second part of the output is a stream of responses:

```

RESPONSE
update {
  timestamp: 1567895852136043891
  update {
    path {
      elem {
        name: "interfaces"
      }
      elem {
        name: "interface"
        key {
          key: "name"
          value: "leaf1-eth3"
        }
      }
    }
    elem {
      name: "state"
    }
    elem {
      name: "counters"
    }
    elem {
      name: "in-unicast-pkts"
    }
  }
  val {
    uint_val: 1592
  }
}
}

```

Each response has a timestamp, path, and new value. Because we are sampling, you should see a new update printed every second. Leave this running, while we generate some traffic.

In another window, open the Mininet CLI and start a ping:

```
$ make mn-cli
```

```
*** Attaching to Mininet CLI...
*** To detach press Ctrl-D (Mininet will keep running)
mininet> h1a ping h1b
```

In the first window, you should see the `uint_val` increase by 1 every second while your ping is still running. (If it's not exactly 1, then there could be other traffic like NDP messages contributing to the increase.)

You can stop the gNMI subscription using `Ctrl-C`.

Finally, we will monitor link events using gNMI's on-change subscriptions.

Start a subscription for the operational status of the first switch's first port:

```
$ util/gnmi-cli --grpc-addr localhost:50001 sub-onchange \
  /interfaces/interface[name=leaf1-eth3]/state/oper-status
```

You should immediately see a response which indicates that port 1 is UP:

```
RESPONSE
update {
  timestamp: 1567896668419430407
  update {
    path {
      elem {
        name: "interfaces"
      }
      elem {
        name: "interface"
        key {
          key: "name"
          value: "leaf1-eth3"
        }
      }
      elem {
        name: "state"
      }
      elem {
        name: "oper-status"
      }
    }
    val {
      string_val: "UP"
    }
  }
}
```

In the shell running the Mininet CLI, let's take down the interface on `leaf1` connected to `h1a`:

```
mininet> sh ifconfig leaf1-eth3 down
```

You should see a response in your gNMI CLI window showing that the interface on `leaf1` connected to `h1a` is `DOWN`:

RESPONSE

```
update {
  timestamp: 1567896891549363399
  update {
    path {
      elem {
        name: "interfaces"
      }
      elem {
        name: "interface"
        key {
          key: "name"
          value: "leaf1-eth3"
        }
      }
    }
    elem {
      name: "state"
    }
    elem {
      name: "oper-status"
    }
  }
  val {
    string_val: "DOWN"
  }
}
```

We can bring back the interface using the following Mininet command:

```
mininet> sh ifconfig leaf1-eth3 up
```

You should see another response in your gNMI CLI window that indicates the interface is `UP`.

Extra credit: We can also use gNMI to disable or enable an interface.

Leave your gNMI subscription for operational status changes running.

In the Mininet CLI, start a ping between two hosts.

```
mininet> h1a ping h1b
```

You should see replies being showed in the Mininet CLI.

In a third window, we will use the gNMI CLI to change the configuration value of the `enabled` leaf from `true` to `false`.

```
$ util/gnmi-cli --grpc-addr localhost:50001 set \  
  /interfaces/interface[name=leaf1-eth3]/config/enabled \  
  --bool-val false
```

In the gNMI set window, you should see request indicating the new value for the `enabled` leaf:

```
REQUEST  
update {  
  path {  
    elem {  
      name: "interfaces"  
    }  
    elem {  
      name: "interface"  
      key {  
        key: "name"  
        value: "leaf1-eth3"  
      }  
    }  
  }  
  elem {  
    name: "config"  
  }  
  elem {  
    name: "enabled"  
  }  
}  
val {  
  bool_val: false  
}  
}
```

In the gNMI subscription window, you should see a new response indicating that the operational status of `leaf1-eth3` is `DOWN`:

```
RESPONSE  
update {  
  timestamp: 1567896891549363399  
  update {  
    path {  
      elem {  
        name: "interfaces"
```

```

    }
    elem {
      name: "interface"
      key {
        key: "name"
        value: "leaf1-eth3"
      }
    }
    elem {
      name: "state"
    }
    elem {
      name: "oper-status"
    }
  }
  val {
    string_val: "DOWN"
  }
}
}

```

And in the Mininet CLI window, you should observe that the ping has stopped working.

Next, we can re-nable the port:

```

$ util/gnmi-cli --grpc-addr localhost:50001 set \
  /interfaces/interface[name=leaf1-eth3]/config/enabled \
  --bool-val true

```

You should see another update in the gNMI subscription window indicating the interface is UP, and the ping should resume in the Mininet CLI window.

Congratulations!

Now you know a little about YANG!